

AI Act Compliance Assistant

Objective

Build a multi-agent application that analyzes **one AI use case at a time** using uploaded use-case documents and a built-in **EU AI Act reference corpus**.

The system should produce a grounded first-pass assessment with citations, identify missing information and uncertainties, and support follow-up questions through chat.

Challenge Overview

Users should be able to create an analysis session for one AI use case and upload multiple supporting documents describing it, such as case studies, vendor white papers, technical overviews, product descriptions, policy documents, or internal process notes. The system should synthesize these inputs into a single case understanding.

System should:

- summarize the use case
- extract relevant facts
- retrieve relevant EU AI Act material
- generate an initial assessment
- show citations and supporting evidence

Users should not need to upload the EU AI Act or related guidance themselves. These should already be available in the application as a built-in reference corpus.

Scope

The core analysis should cover:

- whether the described technology appears to qualify as an AI system
- possible EU AI Act risk classification, including sector- or product-specific AI Act considerations where relevant
- prohibited-practice, transparency, labelling, and GPAI obligations, where relevant
- possible roles, such as provider or deployer
- missing information and uncertainties
- practical compliance-related requirements and governance measures under the EU AI Act

In this challenge, “governance” means different structures, tools, and processes used to manage AI’s effects and mitigate risks through the lens of what is required by the EU AI Act. These may include, for example, documentation, monitoring mechanisms, risk management processes, and

human oversight mechanisms.

EU data legislation, national implementation, and other adjacent legal frameworks may be included as bonus capabilities, but they should be clearly separated from the core AI Act analysis. Sector- or product-specific considerations may be included where they are relevant within the scope of the AI Act.

Suggested Reference Corpus

Required

- **EU AI Act – official Regulation text**
<https://data.europa.eu/eli/reg/2024/1689/oj>

Recommended

- **Commission guidelines on prohibited AI practices**
<https://digital-strategy.ec.europa.eu/en/library/commission-publishes-guidelines-prohibited-artificial-intelligence-ai-practices-defined-ai-act>
- **Commission guidelines on AI system definition**
<https://digital-strategy.ec.europa.eu/en/library/commission-publishes-guidelines-ai-system-definition-facilitate-first-ai-acts-rules-application>
- **Commission Code of Practice on Marking and Labelling of AI-generated Content**
<https://digital-strategy.ec.europa.eu/en/library/commission-publishes-second-draft-code-practice-marking-and-labelling-ai-generated-content>
- **Implementation documents list**
<https://artificialintelligenceact.eu/implementation-documents/>

Optional

- **General-Purpose AI Code of Practice**
<https://digital-strategy.ec.europa.eu/en/policies/contents-code-gpai>
- **EU AI Act Explorer**
<https://artificialintelligenceact.eu>
- **Finnish implementation context**
Traficom: <https://traficom.fi/fi/tekoalyn-saantely/tietoa-eun-tekoalyasetuksesta>
TEM: <https://tem.fi/tekoalyasetus>
Valtioneuvosto: <https://valtioneuvosto.fi/hanke?tunnus=TEM044:00/2024>

Teams may use other reliable sources, but should clearly label the source type, such as legislation, official guidance, national implementation material, commentary, or uploaded use-case material.

Core Requirements

The system should:

- support one AI use case per session
- allow multiple uploaded documents per use case
- use a built-in EU AI Act-focused corpus
- generate an initial analysis without requiring the user to write their own prompt or query
- cite uploaded and reference sources
- support follow-up chat in the same session
- separate facts, references, assumptions, uncertainties, and system-generated interpretation
- avoid presenting outputs as final legal advice

Expected Output

The analysis should include:

1. **Use-case summary**
A concise summary based on the uploaded documents.
2. **Extracted facts**
Key facts such as purpose, users, affected persons, sector, input data, outputs, automation level, human oversight, deployment context, use of AI-generated content, use of GPAI, and possible impact on people.
3. **Preliminary EU AI Act assessment**
A reasoned first-pass assessment with relevant citations, including whether the system appears to meet the AI Act definition of an AI system, possible risk classification, transparency or labelling obligations, GPAI obligations, and uncertainties.
4. **Governance observations**
Practical notes on, for example, documentation, risk management, transparency, human oversight, monitoring, logging, accountability, and role clarity.
5. **Missing information and follow-up questions**
Gaps that prevent a confident assessment and questions an expert might ask next.
6. **Citations and evidence separation**
A clear distinction between uploaded-document facts, regulatory references, optional sources, assumptions, and system reasoning.

Multi-Agent Requirement

Participants should implement a multi-agent workflow where agents have distinct responsibilities, exchange intermediate outputs, and show some degree of autonomous behavior.

The system should not simply be a fixed linear prompt chain. Agents should be able to make context-dependent decisions, such as:

- deciding which uploaded-document facts are relevant for AI Act analysis
- choosing which reference materials to retrieve
- identifying when more evidence is needed

- flagging contradictions, missing information, or weak evidence
- deciding when to ask follow-up questions
- triggering a critique or review step when the assessment is uncertain
- revising the analysis based on new information from the user

Teams may design their own architecture, but should explain what each agent does, what decisions it can make autonomously, how agents communicate, and how the system avoids unsupported or overconfident conclusions.

Bonus Capabilities

Teams can receive bonus points for useful additions beyond the core requirements, such as:

- detecting when EU data legislation, national implementation rules, or other adjacent legal frameworks may also be relevant
- retrieving newer AI Act implementation updates
- adding Finnish implementation context
- providing a dedicated analysis path for GPAI definitions and obligations
- asking targeted follow-up questions and updating the assessment
- showing confidence or uncertainty levels
- creating role-based outputs
- generating a checklist of next steps and documentation needs
- exporting the analysis as a report
- adding a critique or red-team agent
- flagging source quality and type
- comparing several use cases at portfolio level

Evaluation

Solutions will be evaluated on:

- **Content quality:** correct use-case understanding and identification of relevant AI Act concepts
- **Regulatory grounding:** use of relevant EU AI Act sources, accurate mapping between source material and the assessment, and useful citations
- **Multi-agent design:** meaningful agent roles, autonomous behavior, intermediate outputs, and justified workflow
- **User experience:** usability, clarity, and useful follow-up chat
- **Technical implementation:** document handling, retrieval quality, and session context
- **Transparency of output:** clear distinction between facts, assumptions, system reasoning, uncertainties, and limitations
- **Innovation:** creative and practical approaches to the core workflow, user experience, agent design, evidence handling, or bonus capabilities

Note to Participants

The EU AI Act is new, and its practical interpretation is still developing. There is currently no court-case precedent directly interpreting it.

Your system should be transparent about evidence, uncertainty, assumptions, and reasoning. The goal is to build a decision-support tool for structured review, not a final compliance authority.